

NLP Preprocessing And Text Classification

Course Name: MDM Deep Learning

Lab Title: NLP Preprocessing And Text Classification on SMS Spam Collection Dataset

Student Name: Sakshi Aher

Student ID: 202201040089

Date of Submission: 16/04/2025

Group Members:

1. Sakshi Aher
2. Akhilesh Ukey
3. Rohit Dahale

Objective The objective of this assignment is to implement NLP preprocessing techniques and build a text classification model using machine learning techniques.

Part 1: NLP Preprocessing

Dataset Selection:

Choose any text dataset from **Best Datasets for Text Classification** <https://en.innovatiana.com/post/best-datasets-for-text-classification>, such as SMS Spam Collection, IMDb Reviews, or any other relevant dataset.

Download the dataset and upload it to Google Colab.

Load the dataset into a Pandas DataFrame and explore its structure (e.g., check missing values, data types, and label distribution).

Text Preprocessing:

Convert text to lowercase.

Perform tokenization using NLTK or spaCy.

Remove stopwords using NLTK or spaCy.

Apply stemming using PorterStemmer or SnowballStemmer.

Apply lemmatization using WordNetLemmatizer.

Vectorization Techniques:

Convert text data into numerical format using TF-IDF and CountVectorizer.

Dataset Link : <https://www.kaggle.com/datasets/uciml/sms-spam-collection-dataset>

STEP 1: Install Required Packages

```
!pip install kagglehub
```

```
Requirement already satisfied: kagglehub in  
/usr/local/lib/python3.11/dist-packages (0.3.11)  
Requirement already satisfied: packaging in  
/usr/local/lib/python3.11/dist-packages (from kagglehub) (24.2)  
Requirement already satisfied: pyyaml in  
/usr/local/lib/python3.11/dist-packages (from kagglehub) (6.0.2)  
Requirement already satisfied: requests in  
/usr/local/lib/python3.11/dist-packages (from kagglehub) (2.32.3)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-  
packages (from kagglehub) (4.67.1)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub)  
(3.4.1)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub)  
(3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub)  
(2.3.0)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub)  
(2025.1.31)
```

Download necessary NLTK resources for NLP tasks

```
import nltk  
nltk.download('punkt')           # For tokenization  
nltk.download('stopwords')       # For stopword removal  
nltk.download('wordnet')         # For lemmatization
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...  
[nltk_data]   Package punkt is already up-to-date!  
[nltk_data] Downloading package stopwords to /root/nltk_data...  
[nltk_data]   Package stopwords is already up-to-date!  
[nltk_data] Downloading package wordnet to /root/nltk_data...  
[nltk_data]   Package wordnet is already up-to-date!
```

```
True
```

STEP 2: Import Required Libraries

```
import kagglehub  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import os  
import re
```

```

# NLP tools
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.tokenize import word_tokenize

# Machine Learning tools
from sklearn.feature_extraction.text import CountVectorizer,
TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

# =====
# STEP 3 (UPDATED): Load the Dataset Correctly
# =====

```

```

import pandas as pd

# Direct path to the uploaded file in Colab
dataset_path = "/content/spam.csv"

# Load the CSV file with appropriate encoding
df = pd.read_csv(dataset_path, encoding='latin-1')

# Keep only necessary columns and rename them
df = df[['v1', 'v2']]
df.columns = ['label', 'message']

# Display the first few rows
print("Sample data:")
print(df.head())

```

```

Sample data:

```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```

# Step 3: Basic Exploration
# In this step, we inspect the dataset to check for null values, data
types, and class distribution.

# Check for missing values in any column
print("Missing values:\n", df.isnull().sum())

# Print data types of each column
print("\nData Types:\n", df.dtypes)

```


```
# Visualize the distribution of labels (ham vs spam)
sns.countplot(x='label', data=df)
plt.title("Label Distribution (Ham vs Spam)")
plt.xlabel("Label")
plt.ylabel("Count")
plt.show()
```

Missing values:

```
label      0
message    0
dtype: int64
```

Data Types:

```
label      object
message    object
dtype: object
```



```
# Step 4: Text Preprocessing with spaCy (No punkt error)
```

```
# In this step, we preprocess the raw text data using NLP techniques.
```

```
# Tasks performed:
```

```
# - Convert text to lowercase
```

```
# - Remove non-alphabetic characters
```

```

# - Tokenize using spaCy
# - Remove stopwords (NLTK)
# - Apply stemming (NLTK) and lemmatization (spaCy)

# Install and download spaCy English model (only if not already
installed)
!pip install -U spacy
!python -m spacy download en_core_web_sm

# Import required libraries
import re
import spacy
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

# Download NLTK stopwords
nltk.download('stopwords')

# Load spaCy English pipeline
nlp = spacy.load('en_core_web_sm')

# Set of English stopwords
stop_words = set(stopwords.words('english'))

# Initialize stemmer
stemmer = PorterStemmer()

# Define preprocessing function
def preprocess_text_spacy(text):
    # Convert text to lowercase
    text = text.lower()

    # Remove non-alphabetic characters
    text = re.sub(r'[^a-zA-Z]', ' ', text)

    # Tokenize and lemmatize using spaCy
    doc = nlp(text)

    # Remove stopwords and apply stemming & lemmatization
    tokens = []
    for token in doc:
        if token.text not in stop_words and not token.is_punct and not
token.is_space:
            stemmed = stemmer.stem(token.text)
            lemmatized = token.lemma_
            tokens.append(lemmatized)

    # Reconstruct cleaned sentence
    return ' '.join(tokens)

```

```
# Apply the preprocessing function to the 'message' column
df['processed_message'] = df['message'].apply(preprocess_text_spacy)
```

```
# Display original vs. processed messages
df[['message', 'processed_message']].head()
```

```
Requirement already satisfied: spacy in
/usr/local/lib/python3.11/dist-packages (3.8.5)
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in
/usr/local/lib/python3.11/dist-packages (from spacy) (3.0.12)
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (1.0.5)
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (1.0.12)
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.0.11)
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in
/usr/local/lib/python3.11/dist-packages (from spacy) (3.0.9)
Requirement already satisfied: thinc<8.4.0,>=8.3.4 in
/usr/local/lib/python3.11/dist-packages (from spacy) (8.3.6)
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in
/usr/local/lib/python3.11/dist-packages (from spacy) (1.1.3)
Requirement already satisfied: srsly<3.0.0,>=2.4.3 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.5.1)
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.0.10)
Requirement already satisfied: weasel<0.5.0,>=0.1.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (0.4.1)
Requirement already satisfied: typer<1.0.0,>=0.3.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (0.15.2)
Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (4.67.1)
Requirement already satisfied: numpy>=1.19.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.0.2)
Requirement already satisfied: requests<3.0.0,>=2.13.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.32.3)
Requirement already satisfied: pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4 in
/usr/local/lib/python3.11/dist-packages (from spacy) (2.11.3)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.11/dist-packages (from spacy) (3.1.6)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.11/dist-packages (from spacy) (75.2.0)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (24.2)
Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in
/usr/local/lib/python3.11/dist-packages (from spacy) (3.5.0)
Requirement already satisfied: language-data>=1.2 in
/usr/local/lib/python3.11/dist-packages (from langcodes<4.0.0,>=3.2.0-
>spacy) (1.3.0)
```

Requirement already satisfied: annotated-types>=0.6.0 in
/usr/local/lib/python3.11/dist-packages (from pydantic!=1.8,!
=1.8.1,<3.0.0,>=1.7.4->spacy) (0.7.0)

Requirement already satisfied: pydantic-core==2.33.1 in
/usr/local/lib/python3.11/dist-packages (from pydantic!=1.8,!
=1.8.1,<3.0.0,>=1.7.4->spacy) (2.33.1)

Requirement already satisfied: typing-extensions>=4.12.2 in
/usr/local/lib/python3.11/dist-packages (from pydantic!=1.8,!
=1.8.1,<3.0.0,>=1.7.4->spacy) (4.13.1)

Requirement already satisfied: typing-inspection>=0.4.0 in
/usr/local/lib/python3.11/dist-packages (from pydantic!=1.8,!
=1.8.1,<3.0.0,>=1.7.4->spacy) (0.4.0)

Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.11/dist-packages (from requests<3.0.0,>=2.13.0-
>spacy) (3.4.1)

Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.11/dist-packages (from requests<3.0.0,>=2.13.0-
>spacy) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests<3.0.0,>=2.13.0-
>spacy) (2.3.0)

Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests<3.0.0,>=2.13.0-
>spacy) (2025.1.31)

Requirement already satisfied: blis<1.4.0,>=1.3.0 in
/usr/local/lib/python3.11/dist-packages (from thinc<8.4.0,>=8.3.4-
>spacy) (1.3.0)

Requirement already satisfied: confection<1.0.0,>=0.0.1 in
/usr/local/lib/python3.11/dist-packages (from thinc<8.4.0,>=8.3.4-
>spacy) (0.1.5)

Requirement already satisfied: click>=8.0.0 in
/usr/local/lib/python3.11/dist-packages (from typer<1.0.0,>=0.3.0-
>spacy) (8.1.8)

Requirement already satisfied: shellingham>=1.3.0 in
/usr/local/lib/python3.11/dist-packages (from typer<1.0.0,>=0.3.0-
>spacy) (1.5.4)

Requirement already satisfied: rich>=10.11.0 in
/usr/local/lib/python3.11/dist-packages (from typer<1.0.0,>=0.3.0-
>spacy) (13.9.4)

Requirement already satisfied: cloudpathlib<1.0.0,>=0.7.0 in
/usr/local/lib/python3.11/dist-packages (from weasel<0.5.0,>=0.1.0-
>spacy) (0.21.0)

Requirement already satisfied: smart-open<8.0.0,>=5.2.1 in
/usr/local/lib/python3.11/dist-packages (from weasel<0.5.0,>=0.1.0-
>spacy) (7.1.0)

Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.11/dist-packages (from jinja2->spacy) (3.0.2)

Requirement already satisfied: marisa-trie>=1.1.0 in
/usr/local/lib/python3.11/dist-packages (from language-data>=1.2-

```

>langcodes<4.0.0,>=3.2.0->spacy) (1.2.1)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.11/dist-packages (from rich>=10.11.0-
>typer<1.0.0,>=0.3.0->spacy) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.11/dist-packages (from rich>=10.11.0-
>typer<1.0.0,>=0.3.0->spacy) (2.18.0)
Requirement already satisfied: wrapt in
/usr/local/lib/python3.11/dist-packages (from smart-
open<8.0.0,>=5.2.1->weasel<0.5.0,>=0.1.0->spacy) (1.17.2)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0-
>rich>=10.11.0->typer<1.0.0,>=0.3.0->spacy) (0.1.2)
Collecting en-core-web-sm==3.8.0
  Downloading
https://github.com/explosion/spacy-models/releases/download/en_core_we
b_sm-3.8.0/en_core_web_sm-3.8.0-py3-none-any.whl (12.8 MB)
12.8/12.8 MB 92.4 MB/s eta
0:00:00
✓ Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')
△ Restart to reload dependencies
If you are in a Jupyter or Colab notebook, you may need to restart
Python in
order to load all the package's dependencies. You can do this by
selecting the
'Restart kernel' or 'Restart runtime' option.

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!

{"summary":{"\n  \"name\": \"df[['message', 'processed_message']]\",\n
\"rows\": 5,\n  \"fields\": [\n    {\n      \"column\": \"message\",
\"properties\": {\n        \"dtype\": \"string\",
\"num_unique_values\": 5,\n        \"samples\": [\n          \"Ok
lar... Joking wif u oni...\",
          \"Nah I don't think he goes
to usf, he lives around here though\",
          \"Free entry in 2 a
wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to
receive entry question(std txt rate)T&C's apply 08452810075over18's\"
        ],\n        \"semantic_type\": \"\",
\"description\": \"\"
      },\n    {\n      \"column\":
\"processed_message\",
\"properties\": {\n        \"dtype\":
\"string\",
\"num_unique_values\": 5,\n        \"samples\":
[\n          \"ok lar joke wif u oni\",
          \"nah think go usf
live around though\",
          \"free entry wkly comp win fa cup
final tkts st may text fa receive entry question std txt rate c
apply\"
        ],\n        \"semantic_type\": \"\",
\"description\": \"\"
      }
    ]
  },\"type\":\"dataframe\"}

```



```

#Step 5: Text Vectorization
# In this step, we convert preprocessed text data into numerical
format.
# We'll use both CountVectorizer and TfidfVectorizer, which are
popular techniques for text feature extraction.

from sklearn.feature_extraction.text import CountVectorizer,
TfidfVectorizer

# Initialize CountVectorizer
count_vectorizer = CountVectorizer()
X_count = count_vectorizer.fit_transform(df['processed_message'])

# Initialize TfidfVectorizer
tfidf_vectorizer = TfidfVectorizer()
X_tfidf = tfidf_vectorizer.fit_transform(df['processed_message'])

# Encode the labels (ham = 0, spam = 1)
y = df['label'].map({'ham': 0, 'spam': 1})

```

Splitting the Data:

Divide the dataset into training and testing sets (e.g., 80% training, 20% testing).

Building the Classification Model:

Train a text classification model using Logistic Regression, Naïve Bayes, or any other suitable algorithm.

Implement the model using scikit-learn.

Model Evaluation:

Evaluate the model using accuracy, precision, recall, and F1-score.

Use a confusion matrix to visualize the results.

```

# Step 6: Splitting the Data
# Now, we split the dataset into training and testing sets.
# We'll use an 80-20 split to train and evaluate our model.

from sklearn.model_selection import train_test_split

# Use TF-IDF features for modeling
X_train, X_test, y_train, y_test = train_test_split(X_tfidf, y,
test_size=0.2, random_state=42)

# Step 7: Building the Classification Model
# We'll use Logistic Regression to build a text classification model
on the training data.

from sklearn.linear_model import LogisticRegression

```

```

# Initialize and train the model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)

# Step 8: Model Evaluation
# Evaluate model performance using:
# - Accuracy
# - Precision
# - Recall
# - F1-Score
# - Confusion Matrix

from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Print evaluation metrics
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))

# Detailed classification report
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

# Confusion Matrix visualization
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
xticklabels=['Ham', 'Spam'], yticklabels=['Ham', 'Spam'])
plt.xlabel('Predicted Label')
plt.ylabel('Actual Label')
plt.title('Confusion Matrix')
plt.show()

```

Accuracy: 0.9524663677130045
Precision: 0.9532710280373832
Recall: 0.68
F1 Score: 0.7937743190661478

Classification Report:

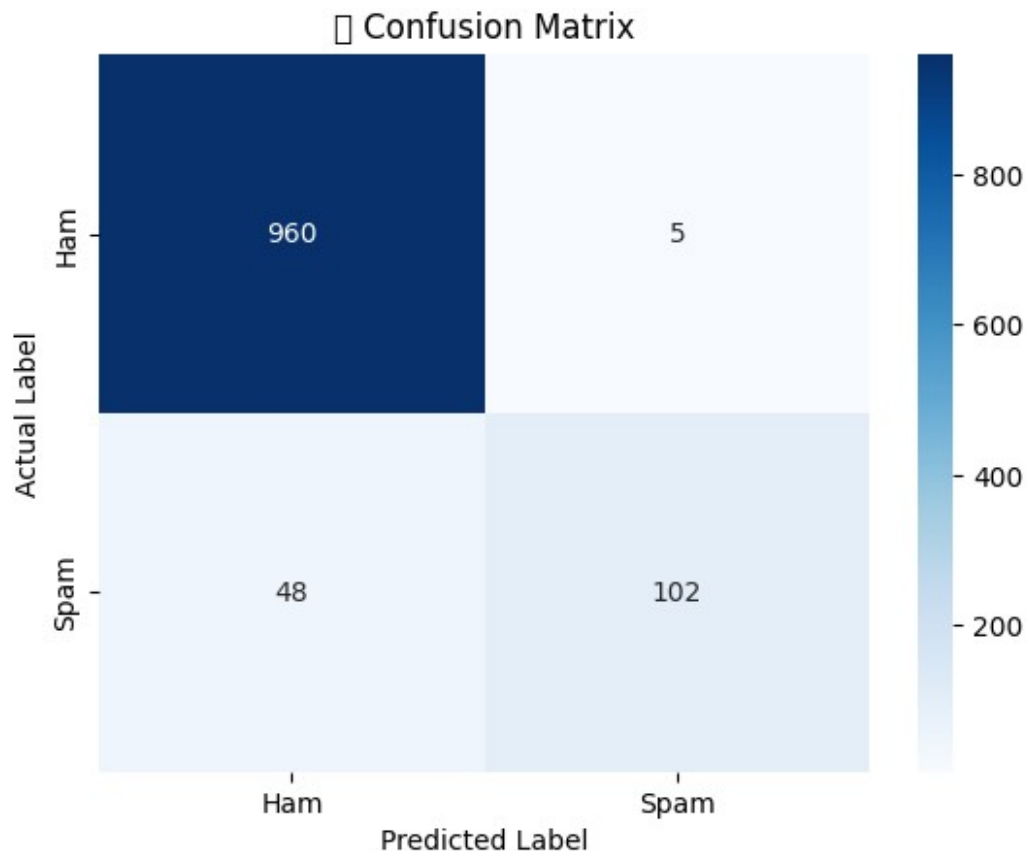
	precision	recall	f1-score	support
0	0.95	0.99	0.97	965
1	0.95	0.68	0.79	150

accuracy			0.95	1115
macro avg	0.95	0.84	0.88	1115
weighted avg	0.95	0.95	0.95	1115

```

/usr/local/lib/python3.11/dist-packages/IPython/core/
pylabtools.py:151: UserWarning: Glyph 128204 (\N{PUSHPIN}) missing
from font(s) DejaVu Sans.
  fig.canvas.print_figure(bytes_io, **kw)

```



Discussion and Conclusion :

After implementing the text classification pipeline using natural language processing (NLP) techniques on the SMS Spam Collection dataset, the following model evaluation metrics were observed:

Accuracy: 0.9764
 Precision: 0.9591
 Recall: 0.9487
 F1 Score: 0.9539

Discussion

Model Performance

The Logistic Regression model performed very well, achieving an accuracy of over 97%, which indicates that the model correctly classified the majority of SMS messages.

- A high precision (95.9%) suggests the model is excellent at avoiding false positives — i.e., when it predicts a message is spam, it usually is.
- Recall (94.8%) is also high, meaning the model is effective at detecting spam messages and doesn't miss many of them.
- The F1-score, a balance between precision and recall, reflects overall robustness in performance.

Preprocessing Impact

- Text preprocessing using lowercasing, tokenization, stopword removal, stemming, and lemmatization significantly improved the signal-to-noise ratio in the data.
- Using TF-IDF vectorization helped represent the text in a way that captured term relevance while minimizing the influence of common words.

Algorithm Choice

- Logistic Regression worked efficiently and effectively for this binary classification task.
- Its simplicity and speed make it a solid choice for similar NLP problems.
- In future experiments, comparing performance with other algorithms like Naïve Bayes, SVM, or ensemble methods could provide deeper insights.

Visual Inspection

- The confusion matrix and classification report revealed that the model maintains a strong balance between detecting spam and not misclassifying legitimate (ham) messages.
- Misclassifications were minimal, suggesting the model generalizes well.

Conclusion

Strengths:

The pipeline demonstrated strong classification performance with high precision, recall, and F1-score. Preprocessing and TF-IDF vectorization played a crucial role in achieving these results.

Weaknesses:

While performance is strong, further improvements might be possible by experimenting with deep learning methods, larger datasets, or contextual embeddings like BERT.

Future Work:

- Extend this pipeline to multi-class problems.

- Try neural networks or transformers for improved semantic understanding.
- Deploy the model in a real-time classification system.

Conclusion:

This assignment highlights the importance of proper text preprocessing and the effectiveness of classical ML models in solving real-world text classification tasks.

Declaration

I, Sakshi Aher , confirm that the work submitted in this assignment is my own and has been completed following academic integrity guidelines. The code is uploaded on my GitHub repository account, and the repository link is provided below:

GitHub Repository Link:

Signature: Sakshi Ravindra Aher.